

# Julia 1: Identifiers, simple types, basic operations

Bernt Lie

University of South-Eastern Norway

May, 2021; minor update September 2023

## Learning goals

In this session, we will look at:

- valid Julia identifiers + naming conventions
- give an overview of simple types
- give an overview of basic functions

## Julia identifiers

### What is an identifier?

By *identifier* (label) is meant a symbolic name for items/memory cells that are introduced by Julia or the user. Some general conventions in computer science:

- **CamelCase** (= **CapitalizedWord** = **CapWords**) where each word starts with a capital letter
- **mixedCase** where the first word starts with a lower case letter while each subsequent word starts with a capital letter
- **underscore\_separated\_words** where each word is separated by underscore
- **combinedwords** where words are simply concatenated

### Valid symbols in identifiers

Variable names must begin with a letter (A-Z or a-z), underscore ("\_"), or a subset of Unicode code points greater than 00A0. [A code point is a number assigned to represent an abstract character in a system for representing text (such as Unicode). In Unicode, a code point is expressed in the form "U+1234" where "1234" is the assigned number. For example, the character "A" is assigned a code point of U+0041].

Subsequent characters may also include ! and digits (0-9 and some other characters), as well as other Unicode code points: diacritics and other modifying marks (categories Mn/Mc/Me/Sk), some punctuation connectors (category Pc), primes, and a few other characters.

Operators such as + are also valid identifiers, but are parsed specially. In some contexts, operators can be used just like variables; for example (+) refers to the addition function, and (+) = f will reassign it.

*Assignment* is expressed by symbol `=` in Julia. Thus, e.g., `x = 2+3` means that the value (here: 5) on the right-hand-side of symbol `=` is *assigned* to identifier `x` on the left-hand-side of symbol `=`; in other words "take the evaluated expression `2+3`  $\rightarrow$  5 on the right-hand-side of symbol `=` and put this value into the memory location given by identifier `x`".

The only explicitly *illegal* names for variables are the names of the built-in Keywords, which at Julia v.1.9.3 are:

```
PS C:\Users\Bernt_Lie> julia --lisp
;
; | _
; | _ _ _ | _ _ | . _ _
; | ( - | | | _ ( _ ) | _ _ ) | _
; -----|-----
-

> reserved-words
(begin while if for try return break continue function macro quote let local
 global const do struct module baremodule using import export end else elseif
 catch finally true false)
```

Examples:

In [532...

```
x = 5.0
y = 3+2im      # = 3+2*j = 3+2*i, complex number
α = 3.5        # \alpha + TAB gives α; use LaTeX syntax
Γ₁ = -2        # \Gamma + TAB followed by \_1 + TAB
🕒 = 8.00       # \: + TAB gives intelligence pop-up menu
;              # semicolon suppresses output
```

In [533...

```
pi
```

```
π = 3.1415926535897...
```

Observe that `pi` is the mathematical constant  $\pi$ , which can also be written as Unicode character `\pi` + TAB.

Using non-ASCII symbols should be used with care, as some word processors don't support it.

## Identifier conventions in Julia

- **Label** ("variable"): Lower case, word separation by underscore if a combined words sequence is difficult to read. *Examples:* `temperature` , `temperature_indoor` , `Tin` , `T_in`
- **Types and Modules:** *CamelCase* (not *mixedCase*). *Examples:* `Float64` , `Base` , `DifferentialEquations`
- **Functions and macros:** Lower case + avoid underscore. If function combines several actions, underscore may be needed; alternatively: split up into several functions. *Examples:* `println` , `isequal` , `timed` . *Note:* macros are specified with symbol `@` at the start, e.g., macro `benchmark` is invoked by `@benchmark` .

- **Mutating functions:** Functions that modify their arguments *by convention* ends with exclamation mark `!` (sometimes read out loud as "bang"). *Examples:* `plot` is non-mutating (creates a fresh plot), `plot!` ("plot bang") is mutating (modifies an existing plot).

## Unevaluated identifiers

It is possible to tell Julia that the expression on the right-hand-side of an assignment should be assigned to the identifier *without evaluating the right-hand-side* -- this is done by prepending the identifier with `:`. In other words: `:x=2+3` means "assign the unevaluated expression `2+3` to identifier `x`". Unevaluated assignment is used in more advanced methods involving so-called *macros*, and is not discussed much in this introduction.

## Memory location of identifiers

Julia does not provide direct information about the location in computer memory of an identifier. However, utility function `objectid(x)` provides a hexadecimal "expression" of the location, which means that two objects whose identifier has the same `objectid`, will have the same memory location, and thus be identical.

Note also that the memory location of an identifier is really a pointer to the first byte of the stored object. Objects taking more space than one byte may be spread over non-contiguous memory space; the data structure that is stored at the `objectid` location will hold information about where the object is stored in memory, and the user does not need to keep track of this.

However, the `objectid` function gives information about whether identifiers are moved in memory, etc.

## Comments and multiline statement continuation

Anything on a line after character `#` is considered a comment, and is not executed.

Multiline comments are enclosed in `#=` and `=#`.

Any incomplete statement at line break leads to an assumed continuation of the statement to the next line.

In [534...

```
# This is a comment
a = 3 # This is also a comment
```

3

In [535...

```
#=
We can have multiline
comments, too
=#
a = 3
```

3

In [536...

```
# Because each of the lines *at line break* are complete statements, the  
# following code is treated as to separate statements  
a = 3  
+ 3
```

3

In [537...

```
# Because the statement of the first line is incomplete *at line break*,  
# the statement is assumed to continue on the next line  
a = 3 +  
3
```

6

## Simple types

### Checking type

Julia comes with built-in simple types (abstract and concrete types), collections, and the possibility of user-defined types.

Type information can be queried by base function `typeof()` where the argument is the type or an identifier holding the type.

As we will see later, the `typeof()` function also returns information about collection structure. For collections, we can query the *element type* by function `eltype()`. For scalars (as here), the result of `typeof()` and the result of `eltype()` will be the same.

Memory requirement (in Bytes) can be queried by `sizeof()` where the argument is the identifier.

In [538...

```
typeof(x)
```

Float64

In [539...

```
typeof(y)
```

Complex{Int64}

In [540...

```
typeof()
```

Float64

In [541...

```
sizeof(x)
```

8

In [542...

```
sizeof()
```

8

In [543...

```
sizeof(y)
```

16

## Character

A *character* type `Char` is given by a single symbol enclosed in single quotation marks, e.g., `'c'`, `'🔴'`. NOTE: a character is not the same as a string.

In [544...

```
x = 'a'
```

'a': ASCII/Unicode U+0061 (category Ll: Letter, lowercase)

In [545...

```
typeof(x)
```

Char

In [546...

```
y = '🔴'
```

'🔴': Unicode U+23F0 (category So: Symbol, other)

In [547...

```
typeof(y)
```

Char

## Number

### Types and subtypes

A (super-) type is related to a subtype as `subtype <: supertype`, where symbol `<:` means "is subtype of", and in some ways is similar to a subset. Essentially, if some code works for a supertype, it should work for its subtypes.

### Abstract number types

- `Number` is an abstract supertype for all number types
- `Real` is a supertype for all real numbers, and is a subtype of `Number` (`Real <: Number` in Julia syntax, where symbol `<:` means subtype and is in some ways similar to subset)
- `AbstractFloat` is a supertype for all floating point numbers; `AbstractFloat <: Real`.
- `AbstractIrrational` represents an exact irrational value; `AbstractIrrational <: Real`. In computations, an `AbstractIrrational` is rounded to correct precision upon operations with other types.
- `Integer` is a supertype of all integer numbers; `Integer <: Real`.
- `Signed` is a supertype of all signed integer numbers; `Signed <: Integer`.
- `Unsigned` is a supertype of all unsigned integer numbers; `Unsigned <: Integer`.

### Concrete number types

- `Float16` is an IEEE 16 bit floating point number; `Float16 <: AbstractFloat`

- `Float32` is an IEEE 32 bit floating point number; `Float32 <: AbstractFloat`
- `Float64` is an IEEE 64 bit floating point number; `Float64 <: AbstractFloat`
- `BigFloat` is a variable precision floating point number; `BigFloat <: AbstractFloat`
- `Bool` is a boolean number (value: `true` or `false`); `Bool <: Integer`. *Note*: `false` is numerically equal to `0` and `true` is numerically equal to `1`
- `Int8` is a signed 8 bit integer; `Int8 <: Signed`
- `Int16` is a signed 16 bit integer; `Int16 <: Signed`
- `Int32` is a signed 32 bit integer; `Int32 <: Signed`
- `Int64` is a signed 64 bit integer; `Int64 <: Signed`
- `Int128` is a signed 128 bit integer; `Int128 <: Signed`
- `BigInt` is a signed variable precision integer; `BigInt <: Signed`
- `UInt8` is an unsigned 8 bit integer; `UInt8 <: Unsigned`
- `UInt16` is an unsigned 16 bit integer; `UInt16 <: Unsigned`
- `UInt32` is an unsigned 32 bit integer; `UInt32 <: Unsigned`
- `UInt64` is an unsigned 64 bit integer; `UInt64 <: Unsigned`
- `UInt128` is an unsigned 128 bit integer; `UInt128 <: Unsigned`
- `Complex` is a complex number; `Complex <: Number`. The real and imaginary numbers can be of various types.
- `Rational` is a rational number; `Rational <: Real`. The numerator and denominators are of type `Integer`
- `Irrational` is an irrational number; `Irrational <: AbstractIrrational`.

## Creating numbers

- By default, a consecutive sequence of numbers without decimal point becomes an integer word size matching the Operating System/Julia version. Thus, `123` is of type `Int64` by default on a 64 bit computer. System label `Sys.WORD_SIZE` holds the machine word size.
- By default, a consecutive sequence of numbers with one decimal point becomes a floating point number matching the OS/Julia word size, hence `1.23` is of type `Float64` by default on a 64 bit computer.
- In order to increase readability, underscore `_` can be inserted to group digits, e.g., `12_103` is easier to read than `12103`; both numbers have the same value in Julia.
- To specify a certain type for a number, the (Abstract or Concrete) type name operates as an instantiator.
- If possible a number can be converted from the existing type to a new type using the type name as a conversion function.
- Negative numbers (except unsigned numbers) are created by prepending the number with symbol `-`. Prepending *Unsigned* numbers with `-` creates the complement number.

## Some integers

In [548...

```
Sys.WORD_SIZE
```

64

In [549...

```
x = 123
```

123

In [550...

```
typeof(x)
```

Int64

In [551...

```
objectid(x)    # Observe the hex number for the object identification
```

0x1cf1facf3ea7e006

In [552...

```
x = Int8(123)
```

123

In [553...

```
typeof(x)
```

Int8

In [554...

```
objectid(x)    # Observe that by re-assigning a new value to `x`, the new  
                # value is stored in a different memory location than  
                # above
```

0x0000000027a694ee

In [555...

```
x = 1_2_3
```

123

In [556...

```
typeof(x)
```

Int64

In [557...

```
x = BigInt(123)
```

123

In [558...

```
typeof(x)
```

BigInt

In [559...

```
x = false
```

false

In [560...

```
typeof(x)
```

```
Bool
```

```
In [561... x = Bool(1)
```

```
true
```

```
In [562... typeof(x)
```

```
Bool
```

```
In [563... x = -1
```

```
-1
```

```
In [564... typeof(x)
```

```
Int64
```

```
In [565... x = -UInt64(1)
```

```
0xffffffffffffffff
```

```
In [566... typeof(x)
```

```
UInt64
```

## Hexadecimal, octal, binary integers. Unsigned integers

Hexadecimal integers are written with prefix `0x` followed by hexadecimal digits

`0, ..., 9, a, ..., f`. Octal integers are input with prefix `0o` followed by octal digits `0, ..., 7`, but are internally represented as hexadecimal numbers. Binary integers are input with prefix `0b` followed by binary digits `0, 1`, but are internally represented as hexadecimal numbers.

Unsigned integers may be input as hexadecimal numbers.

```
In [567... x = 0x1b
```

```
0x1b
```

```
In [568... Int(x)
```

```
27
```

```
In [569... x = 0o17
```

```
0x0f
```

```
In [570... Int(x)
```

```
15
```



In [571...

```
x = 0b10011
```

0x13

In [572...

```
Int(x)
```

19

In [573...

```
UInt8(19)
```

0x13

## Some real numbers

In [574...

```
x = 1.23
```

1.23

In [575...

```
typeof(x)
```

Float64

In [576...

```
x = Float32(1.23)
```

1.23f0

In [577...

```
x = BigFloat(1.23)
```

1.22999999999999982236431605997495353221893310546875

In [578...

```
typeof(x)
```

BigFloat

In [579...

```
x = 12/13
```

0.9230769230769231

In [580...

```
typeof(x)
```

Float64

## Hexadecimal floating point numbers

Hexadecimal floating point numbers can be written with notation `xpn` where `x` is a hexadecimal mantissa and `n` is an decimal integer; such numbers can be represented as `Float64` .

In [581...

```
x = 0xb.acap-2
```

2.9185791015625

In [582...

```
typeof(x)
```

Float64

## Some rational numbers

In [583...

```
x = 12//13
```

12//13

In [584...

```
typeof(x)
```

Rational{Int64}

In [585...

```
x = Int16(12)//Int16(13)
```

12//13

In [586...

```
typeof(x)
```

Rational{Int16}

In [587...

```
x = Int16(12)//Int32(13)
```

12//13

In [588...

```
typeof(x)
```

Rational{Int32}

In [589...

```
x = Rational(12,13)
```

12//13

## Some complex numbers

In [590...

```
x = 12+3im
```

12 + 3im

In [591...

```
typeof(x)
```

Complex{Int64}

In [592...

```
x = Complex(12,3)
```

12 + 3im

## Some number conversion

In [593...

```
x = 123
```

123

In [594...

```
typeof(x)
```

Int64

In [595...

```
x = Int32(x)
```

123

In [596...

```
typeof(x)
```

Int32

In [597...

```
x = BigFloat(1.23)
```

1.22999999999999982236431605997495353221893310546875

In [598...

```
Float64(22//7) # A simple rational approximation of  $\pi$ 
```

3.142857142857143

In [599...

```
x = Float32(x)
```

1.23f0

`x = Int64(x)` # this is not possible: it is necessary to first round a floating point number to a whole number before it can be converted to an integer; see section on functions.

## Floating point numbers in form $x \cdot 10^n$

A number  $x \cdot 10^n$  where  $x$  is an integer or floating point number and  $n$  is an integer, is written as `xen` ( `Float64` ) or `xfn` ( `Float32` ), e.g.,  $1.0 \cdot 10^{-2}$  is written `1.0e-2` , which leads to a `Float64` type, or as `1.0f-2` which leads to a `Float32` type.

**NOTE:** do *not* write  $x \cdot 10^n$  as `x*10^n` ; see section on *Algebraic operations* below and the warning about *type overflow*.

In [600...

```
x = -3.2e2
```

-320.0

In [601...

```
typeof(x)
```

Float64

In [602...

```
x = Float32(x)
```

```
-320.0f0
```

In [603...

```
x = -3.2f2
```

```
-320.0f0
```

## Number properties and special numbers

The largest and smallest numbers that can be represented by a type is given by functions `typemax()` and `typemin()`, where the argument is a type or a number with given type.

The machine precision for *floating point* numbers is the distance between `1.0` in the given type, and the smallest larger number which can be distinguished. The machine precision indicates the possible accuracy with the given type, and is found by function `eps()` (epsilon) where the argument is a type or a number with given type.

For *floating point numbers*, symbols for (plus/minus) infinity as well as not being a number like any other number (including itself), exists.

- **Positive infinity:** `Inf (Float64)`, `Inf32 (Float32)`, `Inf16 (Float16)`
- **Negative infinity:** `-Inf (Float64)`, `-Inf32 (Float32)`, `-Inf16 (Float16)`
- **Not-a-number:** `NaN (Float64)`, `NaN32 (Float32)`, `NaN16 (Float16)`

In [604...

```
# max/min type values  
typemax(Int16)
```

```
32767
```

In [605...

```
typemin(Int16)
```

```
-32768
```

In [606...

```
typemax(UInt16)
```

```
0xffff
```

In [607...

```
Int32(typemax(UInt16))
```

```
65535
```

In [608...

```
typemin(UInt16)
```

```
0x0000
```

In [609...

```
# Floating point machine precision  
eps(Float64)
```

```
2.220446049250313e-16
```

In [610...

```
x = 1.7
```

1.7

In [611...

```
eps(x)
```

2.220446049250313e-16

In [612...

```
eps(Float32)
```

1.1920929f-7

In [613...

```
eps(BigFloat)
```

1.727233711018888925077270372560079914223200072887256277004740694033718360632485e-77

In [614...

```
# Special floating point numbers  
1/Inf
```

0.0

In [615...

```
1/(-Inf)
```

-0.0

In [616...

```
Inf/Inf
```

NaN

In [617...

```
Inf-Inf
```

NaN

## Mathematical operations

The following is a simplified overview.

### Algebraic operations

Infix notation for addition ( + ), subtraction ( - ), multiplication ( \* ), division ( / ), power ( ^ ); grouping of operation by parenthesis in the normal way. [*Infix* implies that a binary operator (say) + is placed *between* the two arguments such as in a+b ; the alternative is a function notation + (a,b) .]

- For *multiplication*, the multiplication operator ( \* ) can be *skipped* between an identifier and a number (just as in mathematics), e.g., with identifier x , 3x is the same as 3\*x
- If the operations are applied to different datatypes, then the numbers are automatically converted to the most general datatype before execution.
- Division ( / ) leads to a floating point number.
- If division between integers is meant to denote integer division, use operator // , which leads to an exact rational number.

- Taking the power between two integers always leads to an integer. **Note:** beware of the possibility of type overflow

In [618...

```
1+3
```

4

In [619...

```
1. + 3
```

4.0

In [620...

```
3-4
```

-1

In [621...

```
3*4.
```

12.0

In [622...

```
12/3
```

4.0

In [623...

```
12//3
```

4//1

In [624...

```
3^4.
```

81.0

In [625...

```
3.0^4
```

81.0

In [626...

```
3^4
```

81

In [627...

```
10^18
```

1000000000000000000

In [628...

```
10^19
```

-8446744073709551616

In [629...

```
# Observe that 10^19 is an integer which is larger than typemax(Int), which
# Leads to the complement number
Float64(typemax(Int))
```

9.223372036854776e18

In [630...

```
# To avoid overflow, use exponent notation!, which produces a
# floating point number.
1e19
```

1.0e19

In [631...

$$1+2/(3-4)$$

-1.0

## Updating operators

- **Addition:**  $x = x + y$  can be written as  $x += y$
- **Subtraction:**  $x = x - y$  can be written as  $x -= y$
- **Multiplication:**  $x = x * y$  can be written as  $x *= y$
- **Division:**  $x = x / y$  can be written as  $x /= y$ ;  $x = x \backslash y$  can be written as  $x \backslash= y$ ;  $x = x \div y$  can be written as  $x \div= y$
- **Power:**  $x = x^y$  can be written as  $x ^= y$

In [632...

$$x = 3$$
$$x = x^2$$

9

In [633...

$$\begin{aligned}x &= 3 \\ x \wedge 2 &= 2\end{aligned}$$

9

## Relational operators

Relational operators are operators on various datatypes which gives a Boolean outcome. Infix notation for equality ( == ), greater than ( > ), greater or equal than ( >= ), less than ( < ), less than or equal ( <= ), not equal to ( != ), identical equal to ( === ), not identical equal to ( !== ); grouping of operation by parenthesis in the normal way.

Because floating point numbers are approximated in a binary representation, we really often will accept numbers as equal if they are *approximatly* equal. Julia provides function `isapproximate()` for that.

In [634...

[illegible]

true

In [635...

[illegible]

```
false
```

```
In [636...
```

```
Float64(1/3) > Float32(1/3)
```

```
false
```

```
In [637...
```

```
Float64(1/3) < Float32(1/3)
```

```
true
```

```
In [638...
```

```
pi ≤ 3.15    # Using LaTeX notation, \le + TAB gives Unicode symbol ≤
```

```
true
```

```
In [639...
```

```
Float64(1/3) != Float32(1/3)
```

```
true
```

```
In [640...
```

```
x = 2  
y = x  
z = 2
```

```
2
```

```
In [641...
```

```
x === y
```

```
true
```

```
In [642...
```

```
x === z
```

```
true
```

```
In [643...
```

```
x !== z
```

```
false
```

```
In [644...
```

```
<(2,3)
```

```
true
```

```
In [645...
```

```
isapprox(Float64(1/3),Float32(1/3))
```

```
true
```

## Logical operators

Logical operators are operators on the `Bool` datatype which gives a Boolean outcome. Infix notation for *negation* ( `!` ), logical *and* ( `&&` ), logical *or* ( `||` ); grouping of operation by parenthesis in the normal way.



Observe that logical *and* and logical *or* are *short-circuiting* in the sense that in `a && b`, `b` is not evaluated if `a` is `false`, while in `a || b`, `b` is not evaluated if `a` is `true`. This short-circuiting helps speed up the execution of the code.

In [646...

```
true == false
```

```
false
```

In [647...

```
(2<3) && (5>4)
```

```
true
```

In [648...

```
(2<3) || (5<4)
```

```
true
```

## Mathematical functions

### Algebraic operations as functions

Infix operators can also be written as functions:

- **Addition:**  $x+y+\dots+z$  can be written as  $+(x,y,\dots,z)$
- **Subtraction:**  $x-y$  can be written as  $-(x,y)$
- **Multiplication:**  $x*y*\dots*z$  can be written as  $*(x,y,\dots,z)$
- **Division:**  $x/y$  can be written as  $/(x,y)$
- **Left division:**  $x\backslash y$  can be written as  $\backslash(x,y)$
- **Power:**  $x^y$  can be written as  $^(x,y)$

In [649...

```
2+5, +(2,5)
```

```
(7, 7)
```

In [650...

```
2/3, /(2,3)
```

```
(0.6666666666666666, 0.6666666666666666)
```

In [651...

```
2^3, ^(2,3)
```

```
(8, 8)
```

In [652...

```
2*3*4, *(2,3,4)
```

```
(24, 24)
```

### Relational operators as functions

Infix relational operators can be written as functions:

- **Equality:**  $x==y$  can be written as  $==(x,y)$  ;  $x!=y$  can be written as  $!=(x,y)$  ;  $x===y$  can be written as  $===(x,y)$  ;  $x!==y$  can be written as  $!==(x,y)$  .
- **Larger than:**  $x>y$  can be written as  $>(x,y)$  ;  $x>=y$  can be written as  $>=(x,y)$  .
- **Less than:**  $x>y$  can be written as  $<(x,y)$  ;  $x<=y$  can be written as  $<:(x,y)$  .

Because computers use finite wordlength binary representation, floating point numbers can not be represented exactly. We are therefore often interested in whether numbers are *approximately* equal, within the given machine precision. Julia provides function `isapprox(x,y)` for that.

In [653...

```
==(Float64(1/3),Float32(1/3))
```

```
false
```

In [654...

```
isapprox(Float64(1/3),Float32(1/3))
```

```
true
```

## Basic mathematical functions: single argument

The following *single argument* mathematical functions are given in the on-line Julia documentation at <https://julialang.org/>:

- **Trigonometric [radians]:** `sin` , `cos` , `tan` , `sec` , `csc` , `cot` represent  $\sin(x)$ ,  $\cos(x)$ ,  $\tan(x)$ ,  $\sec(x) = \frac{1}{\cos(x)}$ ,  $\csc(x) = \frac{1}{\sin(x)}$ ,  $\cot(x) = \frac{1}{\tan(x)}$ , respectively, where  $x$  is given in radians.
- **Trigonometric [degrees]:** `sind` , `cosd` , `tand` , `secd` , `cscd` , `cotd` represent  $\sin(x)$ ,  $\cos(x)$ ,  $\tan(x)$ ,  $\sec(x)$ ,  $\csc(x)$ ,  $\cot(x)$ , respectively, where  $x$  is given in degrees.
- **Inverse trigonometric [radians]:** `asin` , `acos` , `atan` , `asec` , `acsc` , `acot` represent  $\arcsin(x)$ ,  $\arccos(x)$ ,  $\arctan(x)$ ,  $\text{arcsec}(x)$ ,  $\text{arccsc}(x)$ ,  $\text{arccot}(x)$ , respectively, where the result is given in radians.
- **Inverse trigonometric [degrees]:** `asind` , `acosd` , `atand` , `asecd` , `acscd` , `acotd` represent  $\arcsin(x)$ ,  $\arccos(x)$ ,  $\arctan(x)$ ,  $\text{arcsec}(x)$ ,  $\text{arccsc}(x)$ ,  $\text{arccot}(x)$ , respectively, where the result is given in degrees.
- **Special trigonometric [radians]:** `sincos` , `sinpi` , `cospi` , `sincospi` , `sinc` , `cosc` represent tuple  $(\sin(x), \cos(x))$ ,  $(\sin(\pi x), \cos(\pi x))$ , tuple  $(\sin(\pi x), \cos(\pi x))$ ,  $\text{sinc}(x) = \frac{\sin(\pi x)}{\pi x}$ ,  $\text{cosc}(x) = \frac{d}{dx}\text{sinc}(x)$ , respectively.
- **Hyperbolic trigonometric [radians]:** `sinh` , `cosh` , `tanh` , `sech` , `csch` , `coth` represent  $\sinh(x)$ ,  $\cosh(x)$ ,  $\tanh(x)$ ,  $\text{sech}(x) = \frac{1}{\cosh(x)}$ ,  $\text{csch}(x) = \frac{1}{\sinh(x)}$ ,  $\text{coth}(x) = \frac{1}{\tanh(x)}$ , respectively, where  $x$  is given in radians.
- **Inverse hyperbolic [radians]:** `asinh` , `acosh` , `atanh` , `asech` , `acsch` , `acoth` represent  $\text{arsinh}(x)$ ,  $\text{arcosh}(x)$ ,  $\text{artanh}(x)$ ,  $\text{arcsech}(x)$ ,  $\text{arccsch}(x)$ ,  $\text{arcoth}(x)$ , respectively, where  $x$  is given in radians.

- **Angle conversion:** `deg2rad` , `rad2deg` represent  $x \rightarrow \frac{2\pi \cdot x}{360}$  and  $x \rightarrow \frac{360 \cdot x}{2\pi}$ , respectively.
- **Logarithmic/exponential:** `log` , `log2` , `log10` , `log1p` , `exp` , `exp2` , `exp10` , `expm1` represent  $\ln(x) = \log_e(x)$ ,  $\log_2(x)$ ,  $\lg(x) = \log_{10}(x)$ ,  $\ln(1+x)$ ,  $\exp(x)$ ,  $2^x$ ,  $10^x$ , and  $\exp(x) - 1$ , respectively. `log1p` is more accurate than `log` when the argument is close to 1 in value.
- **Roots:** `sqrt` , `cbrt` represent  $\sqrt{x}$  (with an error if  $x$  is real and negative) and  $\sqrt[3]{x}$ , respectively.
- **Absolute values:** `abs` , `abs2` represent  $|x|$  and  $|x|^2$ , respectively. `sign` returns zero if  $x=0$  and  $x/|x|$  otherwise (i.e.,  $\pm 1$  for real  $x$ ).
- **Complex numbers:** with  $j = \sqrt{-1}$  and  $z = x + y \cdot j$ , functions `real` , `imag` , `reim` , `conj` , `angle` , `cis` , `cispi` return  $x$ ,  $y$ , tuple  $(x, y)$ ,  $\bar{z} = x - y \cdot j$ ,  $\arctan\left(\frac{y}{x}\right)$ ,  $\exp(z \cdot j)$ ,  $\exp(2\pi z \cdot j)$ , respectively.
- **Combinatorics:** `factorial(n)` returns  $n!$
- **Rational numbers:** `numerator(x)` returns the numerator of rational number  $x$  , `denominator(x)` returns the denominator of rational number  $x$

In [655...

```
sin(pi/2)
```

1.0

In [656...

```
sind(90)
```

1.0

In [657...

```
sincos(pi/3)
```

(0.8660254037844386, 0.5000000000000001)

In [658...

```
rad2deg(pi/2)
```

90.0

In [659...

```
log10(10)
```

1.0

In [660...

```
cbrt(8)
```

2.0

In [661...

```
abs2(-2)
```

4

In [662...

```
sign(-2), sign(0)
```

(-1, 0)

In [663...

```
real(2+3im)
```

2

In [664...

```
reim(2+3im)
```

(2, 3)

In [665...

```
conj(2+3im)
```

2 - 3im

In [666...

```
angle(2+3im)
```

0.982793723247329

In [667...

```
rad2deg(angle(2+3im))
```

56.309932474020215

In [668...

```
cis(2+3im)
```

-0.02071873100224288 + 0.045271253156092976im

In [669...

```
exp(im*(2+3im))
```

-0.02071873100224288 + 0.045271253156092976im

In [670...

```
cispi(2+3im)
```

8.069951757030463e-5 + 0.0im

In [671...

```
exp(pi*im*(2+3im))
```

8.069951757030463e-5 - 1.976568117704183e-20im

In [672...

```
factorial(1),factorial(2),factorial(3),factorial(4)
```

(1, 2, 6, 24)

In [673...

```
denominator(22//7)
```

7

**Basic mathematical functions: two or more arguments**

The following *multiple argument* mathematical functions are given in the on-line Julia documentation at <https://julialang.org/>:

- **Geometry:** `hypot` [hypotenuse] represents  $(x_1, \dots, x_n) \rightarrow \sqrt{\sum_1^n |x_i|^2}$ , base `b` logarithm of `x`. For a complex number `z`, `hypot(z)` is `hypot(re(z), im(z))`.
- **Logarithm:** `log(b,x)` represents the logarithm of `x` at base `b`, i.e.,  $\log_b(x) = \frac{\log_c(x)}{\log_c(b)}$
- **Modulo, etc.:** `div(x,y)` (or: `÷(x,y)`, `x÷y`) represents `x/y` truncated to an integer; `mod(x,y)` represents `x` modulo `y`; `rem(x,y)` (or: `%(x,y)`, `x%y`) returns the remainder from Euclidian division; `rem2pi(x)` returns `rem(x,pi)`; `mod2pi(x)` returns `mod(x,pi)`.
- **Integer:** `gcd(x1,...,xn)` returns greatest common divisor of arguments, `lcm(x1,...,xn)` returns least common multiple.
- **Combinatorics:** `binomial(n,k)` returns  $\binom{n}{k}$  where `n` and `k` are integers.
- **Min-max:** `min(x, y, ...)` returns the minimum of the arguments; `max(x, y, ...)` returns the maximum of the arguments; `minmax(x, y)` returns the tuple of min and max of the arguments; `clamp(x, lo, hi)` returns the saturation of `x` in the interval `[lo,hi]`.

In [674...

```
hypot(2,3), hypot(2+3im)
```

```
(3.605551275463989, 3.605551275463989)
```

In [675...

```
log(3,8), log2(8)/log2(3)
```

```
(1.892789260714372, 1.8927892607143724)
```

In [676...

```
div(14,3), div(14.,3)
```

```
(4, 4.0)
```

In [677...

```
mod(14,3)
```

```
2
```

In [678...

```
rem(14,3)
```

```
2
```

In [679...

```
div(14,3)*3 + mod(14,3)
```

```
14
```

In [680...

```
gcd(12,36,45)
```

```
3
```

In [681...

```
lcm(12,36,45)
```

180

In [682...

```
binomial(12,12)
```

1

In [683...

```
binomial(12,9)
```

220

In [684...

```
min(-1,0,-2,3)
```

-2

In [685...

```
max(-1,0,-2,3)
```

3

In [686...

```
clamp(1,-1,3)
```

1

In [687...

```
clamp(-2,-1,3)
```

-1

In [688...

```
clamp(4,-1,3)
```

3

## Operation on numbers

### Floating point deconstruction

- **Mantissa:** `significand(x)` extracts the mantissa of `x` in binary representation, with value of the same type as `x` in the interval  $[1, 2)$
- **Exponent:** `exponent(x)` returns the exponent of a normalized floating point number (i.e., with mantissa in  $[1, 2)$ ). The result is the largest integer `y` such that  $2^y \leq \text{abs}(x)$ .

### Rounding

- **Rounding:** `round(x)` rounds `x` to an integer of the same type as `x` ;  
`round(x;digits=val)` where `val` is an integer, returns `x` rounded to `val` digits;  
`round(x;sigdigits=val)` where `val` is an integer returns `x` rounded to `val` significant digits.
- **Ceiling:** `ceil(x)` returns the nearest integral value of the same type as `x` that is greater than or equal to `x` ; `ceil(x;digits=val)` where `val` is an integer, returns the nearest integral

value of the same type as  $x$  that is greater than or equal to  $x$  to  $val$  digits;  
`ceil(x;sigdigits=val)` where  $val$  is an integer, returns the nearest integral value of the same type as  $x$  that is greater than or equal to  $x$  to  $val$  significant digits.

- **Floor:** `floor(x)` returns the nearest integral value of the same type as  $x$  that is less than or equal to  $x$ ; `floor(x;digits=val)` where  $val$  is an integer, returns the nearest integral value of the same type as  $x$  that is less than or equal to  $x$  to  $val$  digits;  
`floor(x;sigdigits=val)` where  $val$  is an integer, returns the nearest integral value of the same type as  $x$  that is less than or equal to  $x$  to  $val$  significant digits.
- **Trunc:** `trunc(x)` returns the nearest integral value of the same type as  $x$  whose *absolute value* is less than or equal to  $x$ ; `trunc(x;digits=val)` where  $val$  is an integer, returns the nearest integral value of the same type as  $x$  whose absolute value is less than or equal to  $x$  to  $val$  digits; `trunc(x;sigdigits=val)` where  $val$  is an integer, returns the nearest integral value of the same type as  $x$  whose absolute value is less than or equal to  $x$  to  $val$  significant digits.

## Rational approximation

- **Rational numbers:** `rationalize(x)` with  $x$  a floating point number returns a rational number approximately equal to  $x$ , `rationalize(x;tol=val)` returns a rational approximation of  $x$  to tolerance given by  $val$ , `rationalize(T,x)` returns a rational number with numerator/denominator given by type  $T$ .

## BigFloat precision

- **BigFloat:** `BigFloat(x;precision=val)` produces a `BigFloat` type of  $x$ , with  $val$  (integer) number of digits in the mantissa (significand). Function `precision(x)` checks the precision of  $x$ .

In [689...

```
x=100pi
```

```
314.1592653589793
```

In [690...

```
s_x = significand(x)
```

```
1.227184630308513
```

In [691...

```
e_x = exponent(x)
```

```
8
```

In [692...

```
s_x*2^e_x
```

```
314.1592653589793
```

In [693...

```
round(x)
```

314.0

In [694...

```
round(x,digits=3)
```

314.159

In [695...

```
round(x,sigdigits=4)
```

314.2

In [696...

```
ceil(x),ceil(-x)
```

(315.0, -314.0)

In [697...

```
floor(x),floor(-x)
```

(314.0, -315.0)

In [698...

```
trunc(x),trunc(-x)
```

(314.0, -314.0)

In [699...

```
ceil(-x)
```

-314.0

In [700...

```
floor(-x)
```

-315.0

In [701...

```
trunc(-x)
```

-314.0

In [702...

```
rationalize(Float64(pi)), rationalize(Float64(pi);tol=1e-3), rationalize(Int16,Float64(  
(165707065//52746197, 201//64, 355//113, 22//7)
```

In [703...

```
x = BigFloat(pi;precision=12)
```

3.1416

In [704...

```
significand(x)
```

1.57080078125